

04 / EVALUATION PLAN

An eight-week gated pilot

Proposed eight-week evaluation after data-access approval. Timing, cohort size, partner ownership and numerical thresholds must be agreed before evaluation begins.

01 Readiness

Weeks 1–2

Approve data scope; verify identity, hosting, retention and prediction cutoffs.

Gate: replay-safe dataset

02 Historical replay

Weeks 3–4

Use time-ordered held-out cases. Compare the full pipeline with simple baselines.

Gate: agreed offline targets

03 Shadow testing

Weeks 5–7

Log predictions and officer review without automated operational action.

Gate: useful, timely leads

04 Joint review

Week 8

Assess benefit, alert burden and failures. Decide whether to extend or stop.

Gate: documented decision

PILOT SCORECARD

MEASURE	WHAT THE PILOT MUST ESTABLISH
Candidate recall	True cell included anywhere in the candidate set.
Top-3 / joint hit rate	Location ranking and location-plus-time success vs baseline.
Delivered lead time	Withdrawal time minus actual investigator availability time.
Investigator workload	False alerts per case/day and time spent reviewing leads.
Data / outcome coverage	Missing links, stale inputs and unobservable cash-out outcomes.

Pre-register thresholds and an alert budget; report uncertainty and all eligible cases, including candidate misses. Historical replay estimates potential lead time. Only prospective testing measures actual delivery and operational usefulness.